

Relational Query Languages

- Relational database systems are expected to be equipped with a *query language* that assist its users to query the database instances.
- **Query Language** is a Language in which user requests information from the database. Eg: **SQL**
- **Query** = “Retrieval Program”
- **Two types of Query Language:**

1. **Procedural Query Language**

2. **Non-Procedural Query Language**



Query Language

Procedural Language

Non-Procedural Language



Procedural vs. Non-Procedural

1. Procedural Query language:

- In **Procedural query language**, user instructs the system to perform a series of operations to produce the desired results.
- User tells what data to be retrieved from database and how to retrieve it.

2. Non-procedural (or Declarative) Query Language:

- In **Non-procedural query language**, user instructs the system to produce the desired result without telling the step by step process.
- User tells what data to be retrieved from database but doesn't tell how to retrieve it.

Two “Pure” Query languages

- Two “Pure” Query languages or Two Mathematical Query Language:

1. Relational Algebra
2. Relational Calculus
 1. Tuple Relational Calculus
 2. Domain Relational Calculus

Relational Algebra vs. Relational Calculus

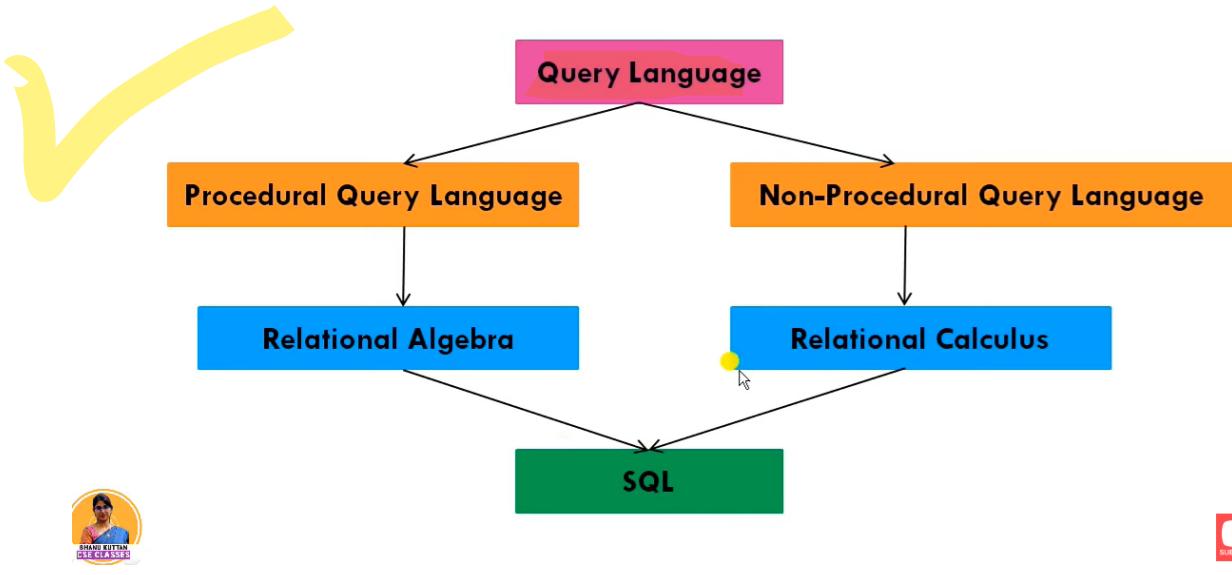
1. Relational Algebra:

- Relational algebra is a procedural query language
- It is more operational, very useful for representing execution plan.
- **Procedural:** What data is required and How to get those data.

2. Relational Calculus:

1. Tuple Relational Calculus
 2. Domain Relational Calculus
- Relational calculus is a non-procedural query language
 - It is non-operational or declarative
 - **Non-Procedural:** What data they want without specifying how to get those data.





Relational Algebra, Calculus, RDBMS & SQL:

- **Relational Model** is a theoretical concept.
- **RDBMS** is a practical implementation of relational model.
- **SQL** (Structured Query Language) is used to write query on RDBMS

- > **Relational algebra and calculus** are the Mathematical system or Query language used on relational model.
 - > Understanding Relational Algebra and Calculus is key to understand **SQL Query Language**
- > **SQL** is a practical implementation of relational algebra and calculus.

Relation Model	RDBMS
Relational Algebra, Relational Calculus	SQL
Algorithm	Code
Conceptual	Reality
Theoretical	Practical



Relational Algebra



- Relational Algebra is a **procedural query language** which takes a relation as an **input** and generates a relation as an **output**
- Relational Algebra is a language for expressing **relational database queries**
- It uses **operators** to perform **queries**. An operator can be either **unary** or **binary**.
- **Types of operators in relational algebra:**
 1. Basic/Fundamental Operators
 2. Additional/Derived Operators
- Relational algebra operations work on one or more relations to define another relation without changing the original relations.
- Relational algebra operations are performed recursively on a relation



Relational Algebra Operations



1. Basic/Fundamental Operations:

1. Selection (σ)
2. Projection (Π)
3. Union (U)
4. Set Difference ($-$)
5. Cartesian product (X)
6. Rename (ρ)

➤ Select, Project and Rename are **Unary** operators because they operate on one relation

➤ Union, Difference and Cartesian product are **Binary** operators because they operate on two relations

2. Additional/Derived Operations:

1. Natural Join (\bowtie)
2. Left, Right, Full Outer Join (\bowtie_l , \bowtie_r , \bowtie_f)
3. Set Intersection (\cap)
4. Division (\div)
5. Assignment (\leftarrow)

➤ All are Binary operators because they operate on two relations





Selection (σ) Operator

- **Selection Operator (σ)** is a **unary operator** in relational algebra that performs a selection operation.
- It selects **tuples (or rows)** that satisfy the **given condition (or predicate)** from a **relation**.
- It is denoted by **sigma (σ)**.
- **Notation =** $\sigma_p(r)$ or $\sigma_{(Condition)}(Relation\ Name)$
 - p is used as a propositional logic formula which may use **logical connectives**: \wedge (AND), \vee (OR), \neg (NOT) and **relational operators** like $=, \neq, <, >, \leq, \geq$ to form the **condition**.
 - The **WHERE clause** of a SQL command corresponds to relational **select $\sigma()$** .
 - **SQL:** SELECT * FROM R WHERE C;
 - **Example:** Select tuples from student table whose age is greater than 17

$\sigma_{age > 17} (Student)$



Examples



- Select tuples from a relation “**Books**” where subject is “**database**”

$\sigma_{subject = "database"} (Books)$

- Select tuples from a relation “**Books**” where subject is “**database**” and price is “**450**”

$\sigma_{subject = "database" \wedge price = 450} (Books)$

- Select tuples from a relation “**Books**” where subject is “**database**” and price is “**450**” or have a publication **year after 2010**

$\sigma_{subject = "database" \wedge price = 450 \vee year > 2010} (Books)$



Projection (Π) Operator

- **Projection Operator (Π)** is a **unary operator** in relational algebra that performs a projection operation.
- It **projects** (or displays) the particular **columns (or attributes)** from a relation and it delete column(s) that are not in the projection list.
- It is denoted by Π
- **Notation –** $\Pi_{A_1, A_2, \dots, A_n}(r)$ or $\Pi_{\text{Attribute_list}}(\text{relation name/table name})$
 - Where A_1, A_2, A_n are attribute names of relation r .
- Duplicate rows are automatically eliminated from result

Student		
roll_no	name	age
1	A	20
2	B	17
3	C	16
4	D	19
5	E	18
6	F	18

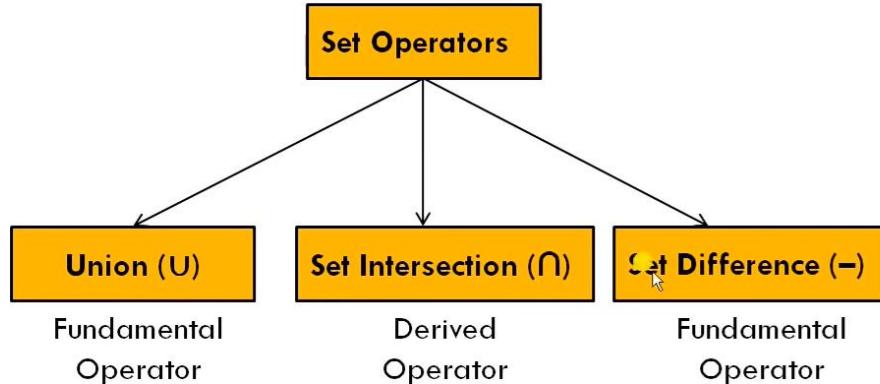
Query 4: Display the roll_no and name of students whose age is greater than 17

$$\Pi_{\text{roll_no, name}} (\sigma_{\text{age} > 17} (\text{Student}))$$



roll_no	name
1	A
4	D
5	E
6	F

Set Operators in Relational Algebra



Set Operators



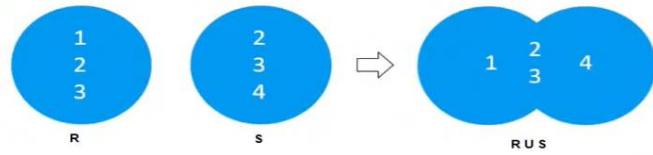
- **Set operators:** Union, intersection and difference, binary operators as they takes two input relations
- **To use set operators on two relations,**
 - The two relations must be **Compatible**
- **Two relations are Compatible if -**
 1. Both the relations must have **same number of attributes (or columns)**.
 2. Corresponding **attribute (or column)** have the **same domain (or type)**.
- Duplicate tuples are automatically eliminated





Union (U) Operator

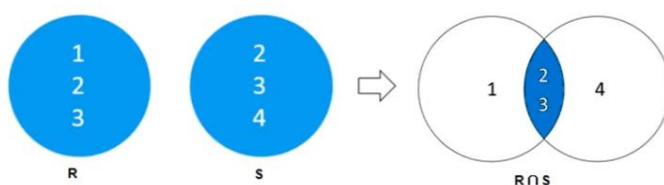
- Suppose R and S are two relations. **The Union operation selects all the tuples that are either in relations R or S or in both relations R & S.**
- It eliminates the **duplicate tuples**.
- For a **union operation** to be **valid**, the following conditions must hold -
 1. Two relations R and S both have **same number of attributes**.
 2. Corresponding **attribute (or column)** have the **same domain (or type)**.
 - The attributes of R and S must occur in the same order.
 3. Duplicate tuples should be automatically removed
- **Symbol:** U
- **Notation:** R U S
 - RA: R U S
 - SQL: (SELECT * FROM R)
UNION
(SELECT * FROM S);



Set Intersection (\cap) Operator



- Suppose R and S are two relations. **The Set Intersection operation selects all the tuples that are in both relations R & S.**
- For a **Set Intersection** to be **valid**, the following conditions must hold –
 1. Two relations R and S both have **same number of attributes**.
 2. Corresponding **attribute (or column)** have the **same domain (or type)**.
 - The attributes of R and S must occur in the same order.
- **Symbol:** \cap
- **Syntax:** R \cap S
 - RA: R \cap S
 - SQL: (SELECT * FROM R)
INTERSECT
(SELECT * FROM S);





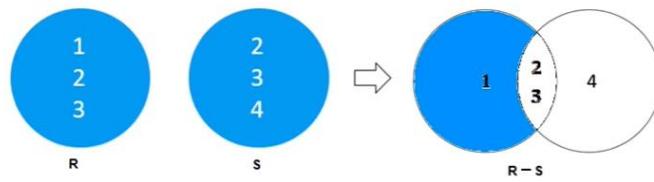
Set Difference (-) Operator

- Suppose R and S are two relations. **The Set Difference operation selects all the tuples that are present in first relation R but not in second relation S.**
- For a **Set Difference** to be **valid**, the following conditions must hold -
 1. Two relations R and S both have **same number of attributes**.
 2. Corresponding **attribute (or column)** have the **same domain (or type)**.
 - The attributes of R and S must occur in the same order.

□ **Symbol:** -

□ **Syntax:** $R - S$

- RA: $R - S$
- SQL: **(SELECT * FROM R)
EXCEPT
(SELECT * FROM S);**



Cartesian Product/Cross Product



- **Cartesian Product** is fundamental operator in relational algebra
- **Cartesian Product combines information of two different relations into one.**

□ It is also called **Cross Product**.

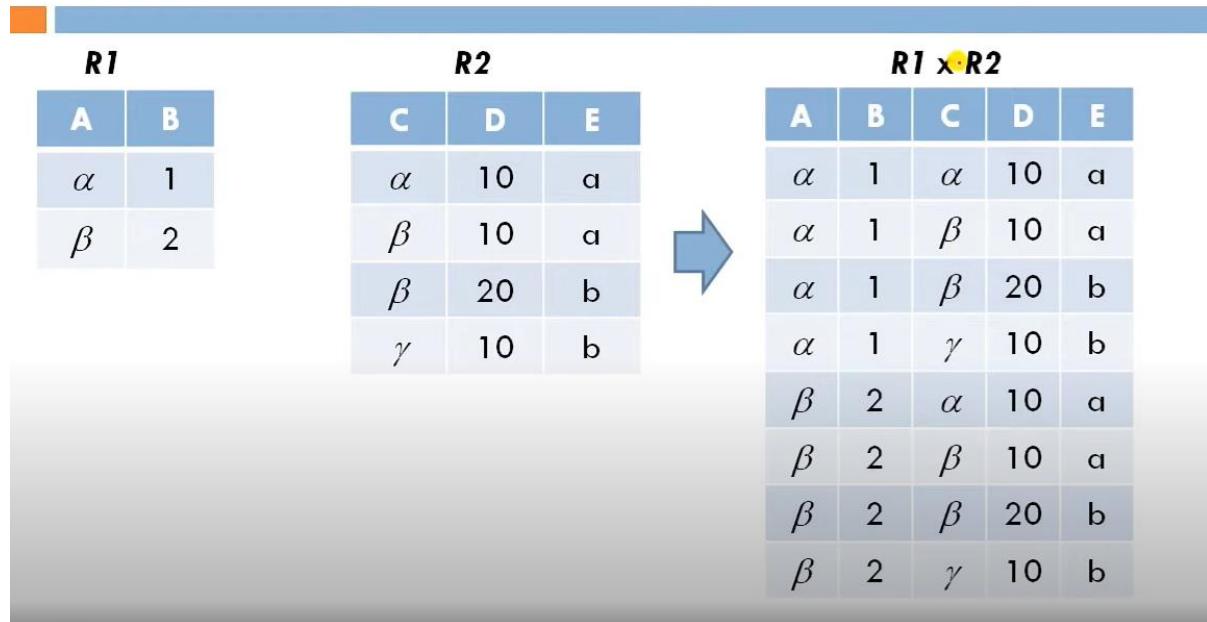
- Generally, a **Cartesian Product** is never a meaningful operation when it is performed alone. However, it becomes **meaningful when it is followed by other operations**.
- Generally it is followed by select operations.

□ **Symbol:** \times

□ **Notation:** $R_1 \times R_2$



Example



Rename Operator



- The results of relational algebra are also relations but without any name.
- **The RENAME operator is used to rename the output of a relation.**
- Sometimes it is simple and suitable to break a complicated sequence of operations and rename it as a relation with different names. Reasons to rename a relation can be many, like:
 - We may want to save the result of a relational algebra expression as a relation so that we can use it later.
 - We may want to join (or cartesian product) a relation with itself, in that case, it becomes too confusing to specify which one of the tables we are talking about, in that case, we rename one of the tables and perform join operations on them.
- **Symbol: rho ρ**
- **Notation 1: $\rho_x(E)$**
Where the symbol ' ρ ' is used to denote the RENAME operator
and **E** is the result of expression or sequence of operation which is saved with the name **X**

- **SQL:** Use the **AS keyword in the FROM clause**

(Eg: **Students AS Students1** renames Students to Students1)

SELECT column_name

FROM tablename AS new_table_name

WHERE condition

- **SQL:** Use the **AS keyword in the SELECT clause to rename attributes (columns)**

(Eg: **RollNo AS SNo** renames RollNo to SNo)

SELECT column_name **AS new_column_name**

FROM tablename

WHERE condition



Example:

- Suppose we want to do cartesian product between same table then **one of the table should be renamed with another name**

RxR

- **R x R**

(Ambiguity will be there)

R	
A	B
α	1
β	2

R.A	R.B	R.A	R.B
α	1	α	1
α	1	β	2
β	2	α	1
β	2	β	2

R x $\rho_s(R)$

- **R x $\rho_s(R)$**

(Rename R to S)

R.A	R.B	S.A	S.B
α	1	α	1
α	1	β	2
β	2	α	1
β	2	β	2



Rename Operation ...



- **Notation 2:** $\rho_x(A_1, A_2, \dots, A_n)(E)$

It returns the result of expression E under the name X , and with the attributes renamed to A_1, A_2, \dots, A_n .

- **Notation 3:** $\rho_{(A_1, A_2, \dots, A_n)}(E)$

It returns the result of expression E with the attributes renamed to A_1, A_2, \dots, A_n .

Examples:



- **Example-2:** Query to rename the attributes Name, Age of table Person to N, A.
 $\rho_{(N, A)}(Person)$
- **Example-3:** Query to rename the table name Project to Work and its attributes to P, Q, R.
 $\rho_{Work(P, Q, R)}(Project)$

- **Example-4:** Query to rename the first attribute of the table Student with attributes A, B, C to P.

$\rho_{(P, B, C)}(Student)$

- **Example-4:** Query to rename the table name Loan to L.

$\rho_L(Loan)$





Introduction

- **Cartesian product** of two relations ($A \times B$), gives us all the possible tuples that are paired together.
 - But it might not be feasible in certain cases to take a **Cartesian product** where we encounter huge relations with **thousands of tuples** having a considerable **large number of attributes**.



Join Operation (\bowtie)



- **Join** is an **Additional / Derived operator** which simplify the queries, but does not add any new power to the basic relational algebra.
- **Join** is a **combination** of a **Cartesian product** followed by a **selection process**.
$$\text{Join} = \text{Cartesian Product} + \text{Selection}$$
- A **Join operation** pairs two tuples from different relations, if and only if a given join condition is satisfied.
- **Symbol:** \bowtie
- $$A \bowtie_c B = \sigma_c(A \times B)$$



Difference



Joins (\bowtie)

- Combination of tuples that satisfy the filtering/matching conditions
- Fewer tuples than cross product, might be able to compute efficiently

R		S	
A	B	B	C
1	a	a	3
2	b	b	4

Cartesian Product /Cross Product/Cross Join(X)

- All possible combination of tuples from the relations
- Huge number of tuples and costly to manage

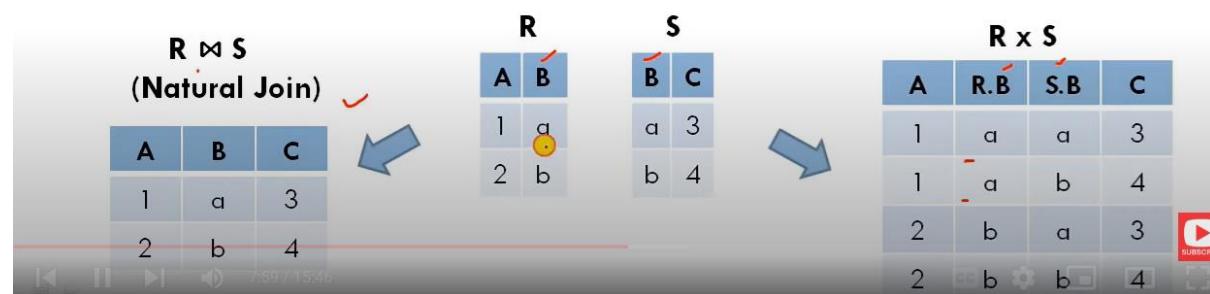


Difference



Joins (\bowtie)

- Combination of tuples that satisfy the filtering/matching conditions
- Fewer tuples than cross product, might be able to compute efficiently



Types of JOINS

1. Inner Join (Join):

- Theta join
- Equi join
- Natural join

2. Outer join:

- (Extension of join)
- Left Outer Join
 - Right Outer Join
 - Full Outer Join

Types of JOINS



1. Inner Join:

- Contains only those tuples that satisfy the matching condition
- Theta(θ) / Conditional Join
 - $A \bowtie_{\theta} B$
 - uses all kinds of comparison operators ($<$, $>$, \leq , \geq , $=$, \neq)
- Equi Join
 - Special case of theta join
 - uses only equality ($=$) comparison operator
- Natural join *MP*
 - $A \bowtie B$
 - Based on common attributes in both relation
 - does not use any comparison operator

2. Outer join:

- Contains matching tuples that satisfy the matching condition, along with some or all tuples that do not satisfy the matching condition
- Contains all rows from either one or both relation
- Left Outer Join
 - Left relation tuples will always be in result whether the value is matched or not
- Right Outer Join
 - Right relation tuples will always be in result whether the value is matched or not
- Full Outer Join
 - Tuples from both relations are present in result, whether the value is matched or not



Inner Join



- An **Inner join** includes **only those tuples that satisfy the matching criteria, while the rest of tuples are excluded.**
- **Theta Join, Equi join, and Natural Join** are called **inner joins.**

1. Theta(θ) / Conditional join



- **Theta join / Conditional Join**
 - It combines **tuples from different relations provided they satisfy the theta (θ) condition.**
 - It is a **general case of join**. And it is **used** when we want to **join** two or more relation based on some **conditions**.
 - The **join condition** is denoted by the symbol θ .
 - It **uses** all kinds of **comparison operators** like $<$, $>$, \leq , \geq , $=$, \neq
- **Notation:** $A \bowtie_{\theta} B$

Where θ is a **predicate/condition**. It can use any comparison operator ($<$, $>$, \leq , \geq , $=$, \neq)
- $A \bowtie_{\theta} B = \sigma_{\theta}(A \times B)$



Example: Theta Join

S1				R1		
sid	name	rating	age	sid	bid	day
22	dustin	7	45.0	22	101	10/10/96
31	lubber	8	55.0	58	103	11/12/96
58	rusty	10	35.0			

$S1 \bowtie_{S1.sid < R1.sid} R1$

S1.sid ✓	sname	rating	age	R1.sid ✓	bid	day
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.0	58	103	11/12/96



2. Equi Join



- When a theta join uses only equivalence ($=$) condition, it becomes a **Equi join**.
- **Equi join** is a special case of theta (or conditional) join where condition contains **equalities** ($=$).
- **Notation:** $A \bowtie_{A.a1 = B.b1 \wedge \dots \wedge A.an = B.bn} B$



Example 1: Equi Join

S1				R1		
sid	name	rating	age	sid	bid	day
22	dustin	7	45.0	22	101	10/10/96
31	lubber	8	55.0	58	103	11/12/96
58	rusty	10	35.0			

$$S1 \bowtie_{S1.sid = R1.sid} R1$$

S1.Sid	sname	rating	age	R1.sid	bid	day
22	dustin	7	45.0	22	101	10/10/96
58	rusty	10	35.0	58	103	11/12/96

Equivalent to $\sigma_{S1.sid = R1.sid}(S1 \times R1)$



3. Natural Join

* jnp



- **Natural join** can only be performed if there is at least one common attribute (column) that exist between two relations. In addition, the attributes must have the same name and domain.
- Natural join does not use any comparison operator.
- It is same as equi join which occurs implicitly by comparing all the common attributes (columns) in both relation, but difference is that in Natural Join the common attributes appears only once. The resulting schema will change.
- **Notation:** A \bowtie B
- The result of the natural join is the set of all combinations of tuples in two relations A and B that are equal on their common attribute names.



Example 1:



$$r = (A, \underline{B}, C, D) \quad s = (\underline{B}, D, E)$$

- Resulting schema of $r \bowtie s = (A, B, C, D, E)$
- $r \bowtie s$ is defined as:

$\prod_{r, A, r, B, r, C, r, D, s, E} (\sigma_{r.B = s.B \wedge r.D = s.D} (r \times s))$

r				s			r \bowtie s				
A	B	C	D	B	D	E	A	B	C	D	E
α	1	α	a	1	a	α	α	1	α	a	α
β	2	γ	a	3	a	β	α	1	α	a	γ
γ	4	β	b	1	a	γ	α	1	γ	a	α
α	1	γ	a	2	b	δ	α	1	γ	a	γ
δ	2	β	B	3	b	ϵ	δ	2	β	b	Δ

Outer Join



- An **Inner join** includes only those tuples with matching attributes and the rest are discarded in the resulting relation. Therefore, we need to use **outer joins** to include all the rest of the tuples from the participating relations in the resulting relation.
- The **outer join** operation is an **extension of the join operation that avoids loss of information**.
- **Outer Join** contains matching tuples that satisfy the matching condition, along with some or all tuples that do not satisfy the matching condition.
 - It is based on both matched or unmatched tuple.
 - It contains all rows from either one or both relations are present
- **It uses NULL values.**
 - NULL signifies that the value is unknown or does not exist



Outer Join Types



- Outer Join = Natural Join + Extra information

(from left table, right table or both table)

- There are three kinds of outer joins:

1. Left outer join
2. Right outer join
3. Full outer join



1. Left outer join ($R1 \bowtie R2$)



When applying **join** on two relations R1 and R2, some tuples of R1 or R2 does not appear in result set which does not satisfy the join conditions. **But..**

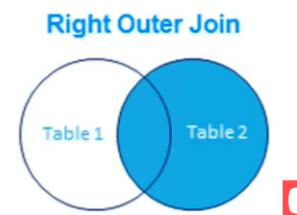
- In **Left outer join**, **all the tuples from the Left relation R1 are included in the resulting relation**. The tuples of R1 which do not satisfy join condition will have values as **NULL** for attributes of R2.
- In short:
 - All record from **left** table
 - Only matching records from right table
- **Symbol:** \bowtie
- **Notation:** $R1 \bowtie R2$



2. Right outer join ($R1 \bowtie R2$)



- In **Right outer join**, **all the tuples from the right relation $R2$ are included in the resulting relation**. The tuples of $R2$ which do not satisfy join condition will have values as **NULL** for attributes of $R1$.
- In short:
 - All record from **right** table
 - Only matching records from left table
- **Symbol:** \bowtie
- **Notation:** $R1 \bowtie R2$



3. Full outer join ($R1 \bowtie R2$)



- In **Full outer join**, **all the tuples from both Left relation $R1$ and right relation $R2$ are included in the resulting relation**. The tuples of both relations $R1$ and $R2$ which do not satisfy join condition, their respective unmatched attributes are made **NULL**.
- In short:
 - All record from **all** table ✓
- **Symbol:** \bowtie
- **Notation:** $R1 \bowtie R2$





Division Operator (\div , /)

- Division operator is a Derived Operator, not supported as a primitive operator
- Suited to queries that include the keyword “all”, or “every” like “at all”, “for all” or “in all”, “at every”, “for every” or “in every”. Eg:
 - Find the person that has account in all the banks of a particular city
 - Find sailors who have reserved all boats.
 - Find employees who works on all projects of company.
 - Find students who have registered for every course.

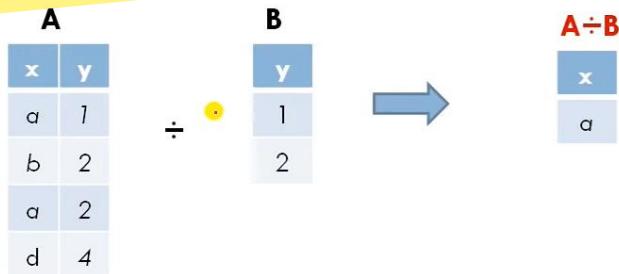


Division Operator (\div)...

Division operator can be applied if and only if:

- Attributes of B is **proper subset** of Attributes of A.
- The relation returned by **division operator** will have **attributes = (All attributes of A – All attributes of B)**.
- The relation returned by **division operator** will **return those tuples from relation A which are associated to every B's tuple.**

Example 1:



Example 2: shows how it works

A ✓		B1 ✓	B2 ✓	B3 ✓
sno	pno	pno	pno	pno
S1	P1	P2	P2	P1
S1	P2		P4	P2
S1	P3			P4
S1	P4			
S2	P1			
S2	P2			
S3	P2			
S4	P2			
S4	P4			

A/B1 ✓	
sno	
S1	
S2	
S3	
S4	

A/B2 ✓	
sno	
S1	✓
S4	✓

A/B3 ✓	
sno	
S1	

Expressing A/B Using Basic Operators



- Division is a derived operator (or additional operator).
- Division can be expressed in terms of **Cross Product, Set-Difference and Projection**.

Idea:

- For A/B, compute all x values that are not 'disqualified' by some y value in B.
- x value is disqualified if by attaching y value from B, we obtain an **xy tuple** that is not in A.

Disqualified x values: $\Pi_X((\Pi_X(A) \times B) - A)$

$$\text{So } A/B = \Pi_X(A) - \text{all disqualified tuples}$$

$$A/B = \Pi_X(A) - \Pi_X((\Pi_X(A) \times B) - A)$$



Expressing A/B Using Basic Operators



- Division is a **derived operator** (or **additional operator**).
- Division can be expressed in terms of **Cross Product, Set-Difference and Projection**.

Idea:

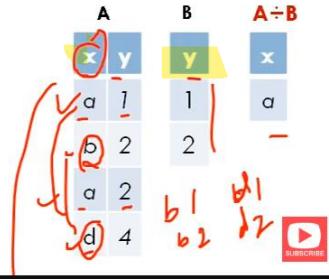
- For A/B , compute all x values that are not 'disqualified' by some y value in B .
 - x value is disqualified if by attaching y value from B , we obtain an **xy tuple** that is not in A .

Disqualified x values:

$$\Pi_X((\Pi_X(A) \times B) - A)$$

So $A/B = \Pi_X(A) - \text{all disqualified tuples}$

$$A/B = \Pi_X(A) - \Pi_X((\Pi_X(A) \times B) - A)$$



Assignment Operator



- The **assignment operation (\leftarrow) provides a convenient way** to express complex queries.
 - It **writes query** as a sequential program consisting of:
 - a series of **assignments**
 - followed by an expression whose value is displayed as a **result** of the query.
 - Assignment** must always be made to a **temporary relation variable**.





Example:

X-
-X

- **Division operation** $A \div B = \Pi_x(A) - \Pi_x((\Pi_x(A) \times B) - A)$
- Write $A \div B$ as

```
temp1 ←  $\Pi_x(A)$ 
temp2 ←  $\Pi_x((temp1 \times B) - A)$ 
result ← temp1 - temp2
```

- The result to the right of the \leftarrow is assigned to the relation variable on the left of the \leftarrow .
- May use variable in subsequent expressions.



Extended Relational Algebra

Extended Relational Algebra increases power over basic relational algebra.

1. Generalized Projection
2. Aggregate Functions
3. Outer Join



1. Generalized Projection

- Normal **projection** only projects the columns whereas **generalized projection** allows arithmetic operations on those projected columns.
- Generalized Projection** extends the **projection operation** by allowing **arithmetic functions** to be used in the **projection list**.

$$\prod_{F_1, F_2, \dots, F_n} (E)$$

- E** is any relational-algebra **expression**
- Each of **F₁, F₂, ..., F_n** are **arithmetic expressions** involving **constants and attributes** in the schema of E.



Example Query

Suggested: Set Operators in Relational Algebra | Union-Intersection...



- Given relation:

`credit_info(customer_name, limit, credit_balance)`

customer_name	limit	credit_balance
A	5000	2000
B	6000	4000
C	10000	6000

"Find how much more each person can spend ?"

$$\prod_{\text{customer_name}, \text{limit} - \text{credit_balance}} (\text{credit_info})$$

customer_name	Limit - credit_balance
A	3000
B	2000
C	4000



Aggregate Functions and Operations



- **Aggregation function** takes a collection of values and returns a single value as a result.

avg: average value
min: minimum value
max: maximum value
sum: sum of values
count: number of values

{
i => 0

- These operations can be applied on **entire relation** or **certain groups of tuples**.
- It ignore **NULL** values except **count**
- Generalize form (g) of **Aggregate operation**:

$G_1, G_2, \dots, G_n \ g_{F_1(A_1), F_2(A_2), \dots, F_n(A_n)} (E)$

- E is any relational-algebra expression
- G_1, G_2, \dots, G_n is a list of attributes on which to group (can be empty)
- Each F_i is an aggregate function
- Each A_i is an attribute name



Aggregate Operation – Example 2



Relation '**account**' grouped by **branch-name**:

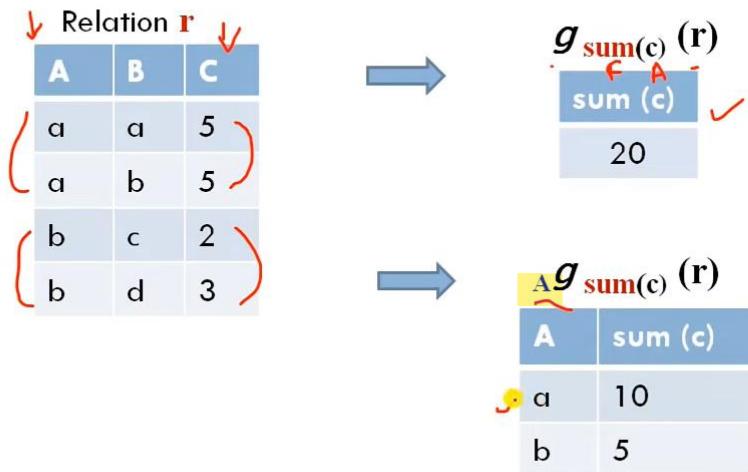
branch_name	account_number	balance
Perryridge	A-102	400
Perryridge	A-201	900
Brighton	A-217	750
Brighton	A-215	750
Redwood	A-222	700

$branch_name \ g \ sum(balance) \ (account)$

branch_name	sum(balance)
Perryridge	1300
Brighton	1500
Redwood	700



Aggregate Operation – Example 1



Modification of the Database



- The **content of the database** may be **modified** using the following operations:
 1. **Deletion**
 2. **Insertion**
 3. **Updating**

1. Deletion

tuples



- A **delete** request is expressed **similarly** to a **query**, except instead of displaying tuples to the user, **the selected tuples are removed from the database**.
- In Deletion, tuples are deleted from the relation ✓ 
- ✓ **Can delete only whole tuples**; cannot delete values on only particular attributes
- A **deletion** is expressed in relational algebra by:

$$r \leftarrow r - E \quad \checkmark$$

where **r** is a relation and **E** is a relational algebra expression.



Example: Deletion

r	A	B	C
✓	1	1	10
✓	1	2	10
✗	2	3	10
✗	2	4	20



A	B	C
1	1	10
1	2	10

$$r \leftarrow r - \{\{1, 2, 10\}\} \quad \checkmark$$

A	B	C
1	1	10
1	2	10

$$r \leftarrow r - \sigma_{A=2}(r) \quad \checkmark$$



Deletion Examples Query



- ✓ Delete all account records in the Perryridge branch.

$\text{account} \leftarrow \text{account} - \sigma_{\text{branch-name} = \text{"Perryridge"}}(\text{account})$

- Delete all loan records with amount in the range of 0 to 50

$\checkmark \text{loan} \leftarrow \text{loan} - \sigma_{\text{amount} \geq 0 \text{ and } \text{amount} \leq 50}(\text{loan})$

- Delete all accounts at branches located in Needham.

$\checkmark r_1 \leftarrow \sigma_{\text{branch-city} = \text{"Needham"}, \text{branch}}(\text{account} \bowtie \text{branch})$

$\checkmark r_2 \leftarrow \prod_{\text{account-number}, \text{branch-name}, \text{balance}}(r_1)$

$r_3 \leftarrow \prod_{\text{customer-name}, \text{account-number}}(r_2 \bowtie \text{depositor})$

$\checkmark \text{account} \leftarrow \text{account} - r_2$

$\text{depositor} \leftarrow \text{depositor} - r_3$

2. Insertion



- Similar to deletion, but uses \cup operator instead of $-$ operator.

- In Insertion, tuples are added to the relation



- To insert data into a relation, we either:

\square specify a tuple to be inserted, or

\square write a query whose result is a set of tuples to be inserted

- An insertion is expressed in relational algebra by:

$r \leftarrow r \cup E$

where r is a relation and E is a relational algebra expression.

- The insertion of a single tuple is expressed by letting E be a constant relation containing one tuple.





3. Updating

- **Updating** is a mechanism to change a value in a tuple without changing all values in the tuple

- Use the generalized projection operator to do this task



- $r \leftarrow \prod_{F_1, F_2, \dots, F_n} (r)$
- Each F_i can be either:
 - the attribute of r , or
 - an expression, involving only constants and the attributes of r , which gives the new value for the attribute
 - **Note:** The schema of the expression resulting from the generalized projection expression must match the original schema of r .



Example: Update



r		
A	B	C
1	1 ✓	10
1	2 ✓	10
2	4 ✓	20

$$r \leftarrow \prod_{A, 2*B, C} (r) \quad \checkmark$$

A	B	C
1	2 ✓	10
1	4 ✓	10
2	8 ✓	20

$$r \leftarrow \prod_{A, 2*B, C} (\sigma_{A=1}(r))$$

A	B	C
1	2 ✓	10
1	4	10



Introduction to Views



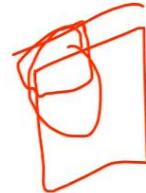
- In some cases, it is not desirable for all users to see the **entire logical model** i.e. **all the actual relations stored in the database.**

■ Consider “a person who needs to know a customer’s loan number but has no need to see the loan amount.” This person should see a relation described, in the relational algebra, by

$$\Pi_{\text{customer-name, loan-number}} (\text{borrower} \bowtie \text{loan})$$

- Any relation that is made visible to a user as a “**virtual relation**” is called a **VIEW**.

- It provides **Limited access to DB.**
- It presents the **Tailored schema.**



What are Views



- “**Views**” are **Virtual Relations** (or **virtual tables**) , through which a **selective portion** of the data from **one or more relations** (or **tables**) **can be seen.**
- Views do not contain data of their own.
- Views do not exist physically.
- **Uses of View:**
 - It helps in **query processing** like **simplify commands** for the user, store complex queries, etc
 - to **restrict access to the database**
 - to **hide data complexity**
- **Database modifications** (like **insert, delete and update operations**) on views **affects the actual relations** in the database, upon which view is based. (Also in SQL)
- Views are stored in the **data dictionary** in the table called **USER_VIEWS**



View Definition



- A **view** is defined using the **create view** statement:

create view v as <query expression>

where **<query expression>** is any legal relational algebra query expression, and **v** represents the **view name**

Example: In SQL,

create view v as

select column-list

from table-name [where condition];

- Once a **view** is defined, the **view name** can be used to refer to the **virtual relation** that the view generates.

Example: In SQL,

select * from v;

- View definition** is not the same as creating a new relation by evaluating the **query expression**.
- Rather, a **view definition causes the saving of an expression to be substituted into queries using the view.**
 - It means wherever **view v** is used, it is actually replaced by the equivalent **query expression** at run time

View in Relational algebra: Example 1



- Creating a view (*loan-customer*) consisting all loan customers and their loan number**

create view loan-customer_as

$\Pi_{customer-name, loan-number} (borrower \bowtie loan)$

- We can find all loan customers and their loan number**

loan-customer

- Note:** So wherever **view *loan-customer*** is used, it is actually replaced by the equivalent **query expression** at run time. This query is evaluated and the entire answer is resulted.



Drop View



- A view can be deleted using the **Drop View** statement.

drop view *viewname*

- Example: **drop view student-view**

Views: Summary



"Views does not stored physically."

- When we define a **view**, the **database system stores the definition of the view itself, rather than the result of evaluation of the relational-algebra expression** that defines the view.
- Wherever a **view relation** appears in a query, it is **replaced by the stored query expression.**
- Thus, whenever we evaluate the query, the **view relation** gets **recomputed.**
- If the original table used in view changes, it does not affects the view relation because it is evaluated every time





Types of Views:

□ **Read-only View :**

- Used only to read data.
- In SQL, it allows only SELECT operations.

□ **Updateable View :**

- Used to read and update the data.
- In SQL, it allows SELECT as well as INSERT , UPDATE and DELETE operations.

Materialized Views



For some views, there is a term called **materialization**. So **some** views are materialized. This depends on the query engine etc., the database engine.

➤ **Certain database systems allow view relations to be stored**, but they make sure that, if the actual relations used in the view definition change, the view is kept up to date. Such views are **called materialized views**.

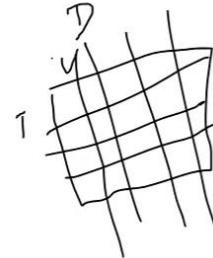
- The **process of keeping the view up to date** is called **view maintenance**.
- Applications that use a view frequently in multiple queries benefited from the use of **materialized views** because then query expression is not going to be evaluated at run time
- Of course, the benefits to queries from the materialization of a view must be weighed against the storage costs and the added overhead for updates.





Relational Calculus

- **Relational Calculus** is a non-procedural query language (or declarative language).
It uses mathematical predicate calculus (or first-order logic) instead of algebra.
- **Relational Calculus** tells what to do but never explains how to do.
- **Relational Calculus** provides description about the query to get the result where as **Relational Algebra** gives the method to get the result.
- When applied to database, it comes in **two flavors**:
 1. **Tuple Relational Calculus (TRC)**:
 - Proposed by Codd in the year 1972
 - Works on tuples (or rows) like SQL
 2. **Domain Relational Calculus (DRC)**:
 - Proposed by Lacroix & pirotte in the year 1977
 - Works on domain of attributes (or Columns) like QBE (Query-By-Example)



Relational Calculus ...

- **Calculus** has variables, constants, comparison operator, logical connectives and quantifiers. $\wedge \vee \neg$
- ✓ **TRC**: **Variables range over tuples**.
 - ✓ Like SQL
- ✓ **DRC**: **Variables range over domain elements**.
 - ✓ Like Query-By-Example (QBE)
- Expressions in the calculus are called **formulas**.
- **Resulting tuple** is an assignment of constants to variables that make the **formula** evaluate to **true**.



1. Tuple Relational Calculus (TRC)



- **Tuple relational calculus** is a non-procedural query language
- **Tuple relational calculus** is used for Selecting the tuples in a relation that satisfy the given condition (or predicate). The result of the relation can have one or more tuples.
- A query in **TRC** is expressed as:

$$\{t | P(t)\}$$

Where t denotes resulting tuple and $P(t)$ denotes predicate (or condition) used to fetch tuple t

- **Result of Query:** It is the set of all tuples t such that predicate P is true for t
- **Notations** used:
 - t is a tuple variable,
 - $t[A]$ denotes the value of tuple t on attribute A
 - $t \in r$ denotes that tuple t is in relation r
 - P is a formula similar to that of the predicate calculus



Predicate Calculus Formula



1. Set of attributes and constants
2. Set of comparison operators: e.g., $<$, \leq , $=$, \neq , $>$, \geq
3. Set of connectives: and (\wedge), or (\vee), not (\neg)
4. Implication (\Rightarrow): $x \Rightarrow y$, if x is true, then y is true $x \Rightarrow y \equiv \neg x \vee y$
5. Quantifiers: **Existential Quantifier (\exists)** and **Universal Quantifier (\forall)**.

- $\exists t \in r (Q(t)) \equiv$ "there exists" a tuple in t in relation r such that predicate $Q(t)$ is true
- $\forall t \in r (Q(t)) \equiv Q$ is true "for all" tuples t in relation r

Free and Bound variables:

- The use of quantifiers $\exists X$ and $\forall X$ in a formula is said to bind X in the formula.
- A variable that is not bound is free.
- Let us revisit the definition of a query: $\{t | P(t)\}$



There is an important restriction

- the variable t that appears to the left of ' $|$ ' must be the only free variable in the formula $P(t)$.
 - in other words, all other tuple variables must be bound using a quantifier





Example Queries: TRC

- Find the loan-number, branch-name, and amount for all loans of over \$1200.

$$\{ t \mid t \in \text{loan} \wedge t[\text{amount}] > 1200 \}$$

It selects all tuples t from relation loan such that the resulting loan tuples will have amount greater than \$1200

- Find the loan number for each loan of an amount greater than \$1200

$$\{ t \mid \exists s \in \text{loan} \ (t[\text{loan-number}] = s[\text{loan-number}] \wedge s[\text{amount}] > 1200) \}$$

It selects the set of tuples t such that there exists a tuple s in relation loan for which the values of t & s for the loan-number attribute are equal and the value of s for the amount attribute is greater than \$1200

branch (branch-name, branch-city, assets)
customer (customer-name, customer-street, customer-city)
account (account-number, branch-name, balance)
loan (loan-number, branch-name, amount)
depositor (customer-name, account-number)
borrower (customer-name, loan-number)



Example Queries: TRC...

- Find the names of all customers having a loan at the Perryridge branch

$$\{ t \mid \exists s \in \text{borrower} \ (t[\text{customer-name}] = s[\text{customer-name}] \wedge \exists u \in \text{loan} \ (u[\text{branch-name}] = \text{"Perryridge"} \wedge u[\text{loan-number}] = s[\text{loan-number}])) \}$$

Join

branch (branch-name, branch-city, assets)
customer (customer-name, customer-street, customer-city)
account (account-number, branch-name, balance)
loan (loan-number, branch-name, amount)
depositor (customer-name, account-number)
borrower (customer-name, loan-number)



- Find the names of all customers having a loan, an account, or both at the bank

$$\{ t \mid \exists s \in \text{borrower} \ (t[\text{customer-name}] = s[\text{customer-name}]) \vee \exists u \in \text{depositor} \ (t[\text{customer-name}] = u[\text{customer-name}]) \}$$

Union

- Find the names of all customers who have a loan and an account at the bank

$$\{ t \mid \exists s \in \text{borrower} \ (t[\text{customer-name}] = s[\text{customer-name}]) \wedge \exists u \in \text{depositor} \ (t[\text{customer-name}] = u[\text{customer-name}]) \}$$

Intersection



2. Domain Relational Calculus (DRC)



- **Domain Relational Calculus** is a non-procedural query language.
- In **Domain Relational Calculus** the records are filtered based on the domains.
- **DRC** uses list of attributes to be selected from relation based on the condition (or predicate).
- **DRC** is same as **TRC** but differs by selecting the attributes rather than selecting whole tuples
- In **DRC**, each query is an expression of the form:
$$\{ \langle a_1, a_2, \dots, a_n \rangle \mid P(a_1, a_2, \dots, a_n) \}$$

a_1, a_2, \dots, a_n represent domain variables

P represents a predicate similar to that of the predicate calculus
- **Result of Query:** It is the set of all tuples $\langle a_1, a_2, \dots, a_n \rangle$ such that predicate P is true for $\langle a_1, a_2, \dots, a_n \rangle$ tuples

Example Queries: DRC



- Find the loan-number, branch-name, and amount for loans of over \$1200:
$$\{ \langle l, b, a \rangle \mid \langle l, b, a \rangle \in \text{loan} \wedge a > 1200 \}$$

(Selection)
- Find the loan number for each loan of an amount greater than \$1200:
$$\{ \langle l \rangle \mid \exists b, a \quad \langle l, b, a \rangle \in \text{loan} \wedge a > 1200 \}$$

(Selection then Projection)



branch (branch-name, branch-city, assets)
customer (customer-name, customer-street, customer-city)
account (account-number, branch-name, balance)
loan (loan-number, branch-name, amount)
depositor (customer-name, account-number)
borrower (customer-name, loan-number)





Example Queries: DRC ...



- Find the names of all customers who have a loan of over \$1200:

$$\{ \underline{\underline{c}} | \exists \underline{\underline{l}}, \underline{\underline{b}}, \underline{\underline{a}} (\underline{\underline{c}}, \underline{\underline{l}}) \in \underline{\underline{\text{borrower}}} \wedge (\underline{\underline{l}}, \underline{\underline{b}}, \underline{\underline{a}}) \in \underline{\underline{\text{loan}}} \wedge \underline{\underline{a}} > 1200) \}$$

7m

- Find the names of all customers having a loan at the Perryridge branch and find the loan amount:

$$\{ \underline{\underline{c}}, \underline{\underline{a}} | \exists \underline{\underline{l}} (\underline{\underline{c}}, \underline{\underline{l}}) \in \underline{\underline{\text{borrower}}} \wedge \exists \underline{\underline{b}} (\underline{\underline{l}}, \underline{\underline{b}}, \underline{\underline{a}}) \in \underline{\underline{\text{loan}}} \wedge \underline{\underline{b}} = \text{"Perryridge"}) \}$$

branch (branch-name, branch-city, assets)
customer (customer-name, customer-street, customer-city)
account (account-number, branch-name, balance)
loan (loan-number, branch-name, amount)
depositor (customer-name, account-number)
borrower (customer-name, loan-number)

